

Enterprise AI Governance Policy

Document Ref: KB-POL-2026-01 | Page 1 of 2

1. Executive Charter & Purpose

The Enterprise AI Governance and Acceptable Use Policy establishes strict operational, security, and ethical boundaries for deploying Large Language Models (LLMs) and autonomous agent workflows within the organization.

This policy applies to all engineering squads, business operations, and automated agent orchestrators operating across production and staging environments.

2. Approved Model Registry & Deployment Tiers

Production agent workflows are authorized to utilize models exclusively from the Enterprise Model Registry. Approved models include Groq LPU-accelerated checkpoints (Llama-3.3-70B, Qwen-2.5-32B), Google Gemini (Gemini-1.5-Flash, Gemini-1.5-Pro), and Claude 3.5 Sonnet.

Experimental or unvetted open-source weights downloaded from public registries without InfoSec audit are strictly prohibited from processing company intellectual property.

3. Data Classification & Strict PII Masking Standards

All automated ingestion pipelines must apply deterministic token masking before vector embedding. Sensitive data categories including credit card numbers (PCI-DSS), Social Security Numbers, protected health information (HIPAA), and customer credentials must never enter vector indexes.

Embeddings generated for vector storage must strictly utilize isolated, authorized tenant namespaces with TLS 1.3 in-flight encryption.

Enterprise AI Governance Policy

Document Ref: KB-POL-2026-01 | Page 2 of 2

4. Grounding & Anti-Hallucination Controls (RAG Architecture)

Retrieval-Augmented Generation (RAG) agents must enforce strict context bounding. System prompts must instruct models to synthesize answers solely from retrieved documentation.

If the retrieval confidence score or similarity threshold falls below acceptable parameters, the agent is mandated to issue a standard insufficient-information disclaimer rather than attempting speculative completion.

5. Human-in-the-Loop Escalation Protocols

Autonomous customer-facing workflows must incorporate automated escalation thresholds. High-sentiment distress, contract termination requests, and high-value refund demands exceeding USD must be immediately routed to certified human operators via Slack or CRM escalation tickets.

6. Compliance Auditing & Log Retention

Every agent inference call, prompt payload, tool invocation, and token utilization metric must be immutably recorded for a minimum of 365 days in compliance with SOC-2 Type II and ISO 27001 audit standards.